

Untitled11

December 15, 2025

1 Time Series Forecasting of NFLX Stock Prices

This project performs exploratory data analysis, stationarity testing, ARIMA modeling, GARCH volatility modeling, and combined ARIMA+GARCH forecasting on Netflix stock prices.

```
[17]: import warnings
warnings.filterwarnings("ignore")

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import matplotlib.dates as mdates
import seaborn as sns
sns.set_style("whitegrid")
%matplotlib inline
```

```
[2]: file_path = r"C:/Users/USER/Music/Self Projects/Time Series forecasting of_
↪Stock Prices/NFLX_stocks.csv"
df = pd.read_csv(file_path)

if 'unnamed: 0' in df.columns:
    df.rename(columns={'unnamed: 0': 'date'}, inplace=True)

df['date'] = pd.to_datetime(df['date'])
df.set_index('date', inplace=True)
df.sort_index(inplace=True)

df['close'] = pd.to_numeric(df['close'], errors='coerce')
df = df.dropna(subset=['close'])

df.head()
```

```
[2]:      open          high         low \
date
2002-05-23  1.1564290523529053  1.2428569793701172  1.1457140445709229
2002-05-24   1.214285969734192   1.225000023841858  1.1971429586410522
2002-05-28   1.2135709524154663   1.2321430444717407  1.157142996788025
2002-05-29   1.1642860174179077   1.1642860174179077  1.0857139825820923
```

```
2002-05-30  1.1078569889068604  1.1078569889068604  1.0714290142059326
```

```
           close      adj_close      volume
date
2002-05-23  1.196429  1.1964290142059326  104790000
2002-05-24  1.210000  1.210000381469727   11104800
2002-05-28  1.157143  1.157142996788025   6609400
2002-05-29  1.103571  1.1035710573196411   6757800
2002-05-30  1.071429  1.0714290142059326  10154200
```

```
[3]: df['close'].describe()
```


```
count      5802.000000
mean       172.807773
std        234.885388
min        0.372857
25%        4.307143
50%        47.187143
75%        317.935005
max       1250.520020
Name: close, dtype: float64
```

```
[4]: plt.figure(figsize=(14,6))
plt.plot(df.index, df['close'], linewidth=1)

plt.title("NFLX Closing Price Time Series")
plt.xlabel("Date")
plt.ylabel("Close Price")

plt.gca().xaxis.set_major_locator(mdates.YearLocator(2))
plt.gca().xaxis.set_major_formatter(mdates.DateFormatter('%Y'))

plt.grid(True, alpha=0.3)
plt.tight_layout()
plt.show()
```



```
[5]: monthly_close = df['close'].resample('M').mean()

plt.figure(figsize=(14,6))
plt.plot(monthly_close, linewidth=2)

plt.title("NFLX Monthly Average Closing Price")
plt.xlabel("Date")
plt.ylabel("Close Price")

plt.gca().xaxis.set_major_locator(mdates.YearLocator(2))
plt.gca().xaxis.set_major_formatter(mdates.DateFormatter('%Y'))

plt.grid(True, alpha=0.3)
plt.tight_layout()
plt.show()
```




```
[6]: plt.figure(figsize=(14,6))
plt.plot(df.index, df['close'], linewidth=1)

plt.yscale('log')
plt.title("NFLX Closing Price (Log Scale)")
plt.xlabel("Date")
plt.ylabel("Log Close Price")

plt.gca().xaxis.set_major_locator(mdates.YearLocator(2))
plt.gca().xaxis.set_major_formatter(mdates.DateFormatter('%Y'))

plt.grid(True, alpha=0.3)
plt.tight_layout()
plt.show()
```




```
[7]: rolling_mean = df['close'].rolling(window=252).mean()
rolling_std = df['close'].rolling(window=252).std()

plt.figure(figsize=(14,6))
plt.plot(df['close'], label='Close Price', alpha=0.6)
plt.plot(rolling_mean, label='252-Day Rolling Mean', linewidth=2)
plt.plot(rolling_std, label='252-Day Rolling Std', linewidth=2)

plt.title("NFLX Rolling Mean & Volatility")
plt.xlabel("Date")
plt.ylabel("Price")

plt.legend()
plt.grid(True, alpha=0.3)
```

```
plt.tight_layout()
plt.show()
```



```
[8]: from statsmodels.tsa.stattools import adfuller, kpss

def adf_test(series):
    return adfuller(series, autolag='AIC')[1]

def kpss_test(series):
    return kpss(series, regression='c', nlags="auto")[1]

print("ADF (level):", adf_test(df['close']))
print("KPSS (level):", kpss_test(df['close']))

df['log_close'] = np.log(df['close'])
df['diff_log_close'] = df['log_close'].diff()

print("ADF (diff log):", adf_test(df['diff_log_close'].dropna()))
print("KPSS (diff log):", kpss_test(df['diff_log_close'].dropna()))
```


```
ADF (level): 1.0
KPSS (level): 0.01
C:\Users\USER\AppData\Local\Temp\ipykernel_53620\1044478671.py:7:
InterpolationWarning: The test statistic is outside of the range of p-values
available in the
look-up table. The actual p-value is smaller than the p-value returned.
```


```
return kpss(series, regression='c', nlags="auto")[1]
ADF (diff log): 0.0
KPSS (diff log): 0.1
```

```
C:\Users\USER\AppData\Local\Temp\ipykernel_53620\1044478671.py:7:  
InterpolationWarning: The test statistic is outside of the range of p-values  
available in the  
look-up table. The actual p-value is greater than the p-value returned.
```

```
    return kpss(series, regression='c', nlags="auto")[1]
```

```
[9]: from statsmodels.graphics.tsaplots import plot_acf, plot_pacf  
  
plot_acf(df['diff_log_close'].dropna(), lags=40)  
plt.show()  
  
plot_pacf(df['diff_log_close'].dropna(), lags=40)  
plt.show()
```






```
[10]: from statsmodels.stats.diagnostic import acorr_ljungbox  
acorr_ljungbox(df['diff_log_close'].dropna(), lags=[10], return_df=True)
```

```
[10]:      lb_stat    lb_pvalue  
10   13.191168   0.213179
```


```
[11]: from statsmodels.tsa.seasonal import seasonal_decompose  
seasonal_decompose(df['log_close'], model='additive', period=252).plot()  
plt.show()
```



```
[12]: df['returns'] = 100 * df['log_close'].diff()

plt.figure(figsize=(14,5))
plt.plot(df['returns'])
plt.title("Log Returns")
plt.show()

from statsmodels.stats.diagnostic import het_arch
het_arch(df['returns'].dropna())
```



```
[12]: (33.15552477022573,
 0.0002564674388313528,
 3.3283103424612555,
 0.00025105222316720497)
```

```
[13]: from arch import arch_model
```

```
garch = arch_model(df['returns'].dropna(),
                   mean='zero',
                   vol='Garch',
                   p=1, q=1)

garch_fit = garch.fit(disp='off')
print(garch_fit.summary())
```

Zero Mean - GARCH Model Results					
Dep. Variable:	returns	R-squared:	0.000		
Mean Model:	Zero Mean	Adj. R-squared:	0.000		
Vol Model:	GARCH	Log-Likelihood:	-15165.7		
Distribution:	Normal	AIC:	30337.4		
Method:	Maximum Likelihood	BIC:	30357.4		
		No. Observations:	5801		
Date:	Mon, Dec 15 2025	Df Residuals:	5801		
Time:	20:48:52	Df Model:	0		
		Volatility Model			
	coef	std err	t	P> t	95.0% Conf. Int.
omega	0.0480	2.761e-02	1.740	8.188e-02	[-6.076e-03, 0.102]
alpha[1]	0.0136	4.255e-03	3.205	1.350e-03	[5.298e-03, 2.198e-02]
beta[1]	0.9825	4.845e-03	202.781	0.000	[0.973, 0.992]

Covariance estimator: robust

```
[14]: from statsmodels.tsa.arima.model import ARIMA

arima = ARIMA(df['log_close'].dropna(), order=(1,1,1))
arima_fit = arima.fit()
print(arima_fit.summary())
```

C:\Users\USER\anaconda3\Lib\site-packages\statsmodels\tsa\base\tsa_model.py:473:
ValueWarning: A date index has been provided, but it has no associated frequency
information and so will be ignored when e.g. forecasting.
 self._init_dates(dates, freq)
C:\Users\USER\anaconda3\Lib\site-packages\statsmodels\tsa\base\tsa_model.py:473:
ValueWarning: A date index has been provided, but it has no associated frequency
information and so will be ignored when e.g. forecasting.
 self._init_dates(dates, freq)
C:\Users\USER\anaconda3\Lib\site-packages\statsmodels\tsa\base\tsa_model.py:473:
ValueWarning: A date index has been provided, but it has no associated frequency
information and so will be ignored when e.g. forecasting.
 self._init_dates(dates, freq)

SARIMAX Results

```
=====
Dep. Variable: log_close    No. Observations: 5802
Model: ARIMA(1, 1, 1)    Log Likelihood: 11202.714
Date: Mon, 15 Dec 2025    AIC: -22399.428
Time: 20:48:53            BIC: -22379.431
Sample: 0                 HQIC: -22392.472
                           - 5802
Covariance Type: opg
=====
```

	coef	std err	z	P> z	[0.025	0.975]
ar.L1	-0.1104	0.362	-0.305	0.760	-0.819	0.599
ma.L1	0.1382	0.360	0.383	0.701	-0.568	0.845
sigma2	0.0012	6.11e-06	201.325	0.000	0.001	0.001

```
=====
```

```
===
Ljung-Box (L1) (Q): 0.00    Jarque-Bera (JB): 169623.70
Prob(Q): 0.95    Prob(JB): 0.00
Heteroskedasticity (H): 0.45    Skew: -1.12
Prob(H) (two-sided): 0.00    Kurtosis: 29.40
=====
```

====

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

```
[15]: # =====
# FIXED ARIMA + GARCH FORECAST
# =====

n_forecast = 30

# ----- ARIMA forecast -----
arima_forecast = arima_fit.get_forecast(steps=n_forecast)
mean_log_forecast = arima_forecast.predicted_mean

# Create proper future date index
last_date = df.index[-1]
forecast_index = pd.date_range(
    start=last_date + pd.Timedelta(days=1),
    periods=n_forecast,
    freq='B' # business days
)

mean_log_forecast.index = forecast_index

# ----- GARCH forecast -----
garch_forecast = garch_fit.forecast(horizon=n_forecast)

# Volatility of RETURNS
volatility = np.sqrt(garch_forecast.variance.iloc[-1].values)

# Scale volatility correctly (returns are %)
volatility = volatility / 100


# ----- Combine correctly -----
forecast_df = pd.DataFrame(index=forecast_index)
forecast_df['mean_log'] = mean_log_forecast.values
forecast_df['upper_log'] = forecast_df['mean_log'] + 1.96 * volatility
forecast_df['lower_log'] = forecast_df['mean_log'] - 1.96 * volatility

# Convert to price
forecast_df['mean_price'] = np.exp(forecast_df['mean_log'])
forecast_df['upper_price'] = np.exp(forecast_df['upper_log'])
forecast_df['lower_price'] = np.exp(forecast_df['lower_log'])
```

C:\Users\USER\anaconda3\Lib\site-packages\statsmodels\tsa\base\tsa_model.py:836:
ValueWarning: No supported index is available. Prediction results will be given

```
with an integer index beginning at `start`.  
    return get_prediction_index()
```

```
[16]: plt.figure(figsize=(14,6))  
  
# Plot recent history only (important)  
plt.plot(df.index[-300:], df['close'][-300:], label='Historical', linewidth=2)  
  
# Plot forecast  
plt.plot(forecast_df.index, forecast_df['mean_price'],  
         label='Forecast', linewidth=2, color='black')  
  
# Confidence bands  
plt.fill_between(  
    forecast_df.index,  
    forecast_df['lower_price'],  
    forecast_df['upper_price'],  
    alpha=0.3,  
    label='95% Confidence Interval'  
)  
  
plt.title("ARIMA + GARCH Forecast (Corrected)")  
plt.xlabel("Date")  
plt.ylabel("Price")  
plt.legend()  
plt.grid(True, alpha=0.3)  
plt.tight_layout()  
plt.show()
```



1.1 Conclusion

Stock prices are non-stationary, while returns are stationary with strong volatility clustering. ARIMA effectively models the conditional mean, and GARCH captures time-varying volatility, making the combined ARIMA + GARCH framework optimal for financial time series forecasting.

[]: